

Name: Ilyas Hakkou

Lab Sqoop

1. Executing the job:

```
[cloudera@quickstart ~]$ sqoop job \
> --create import_orders_inc \
> -- import \
> --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \
> --username retail_dba \
> --password cloudera \
> --table orders_inc \
> --warehouse-dir /user/cloudera/sqoop_import/retail_db \
> --check-column order_id \
> --incremental append \
> --last-value 0
Warning: /usr/lib/sqoop/../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
25/12/04 02:51:18 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
25/12/04 02:51:18 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.
[cloudera@quickstart ~]$
```

2. Inspecting the logs:

```
25/12/04 03:13:12 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN(`order_id`), MAX(`order_id`) FROM `orders_in
c` WHERE ( `order_id` > 0 AND `order_id` <= 30000 )
25/12/04 03:13:12 INFO db.IntegerSplitter: Split size: 7499; Num splits: 4 from: 1 to: 30000
25/12/04 03:13:12 INFO mapreduce.JobSubmitter: number of splits:4
```

We can see that in the request there is a limit, to start from id 0 to 30000.

3. Running job show

```
[cloudera@quickstart ~]$ sqoop job --show import_orders_inc
Warning: /usr/lib/sqoop/../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
25/12/04 03:15:49 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
Enter password:
Job: import_orders_inc
Tool: import
Options:
-----
verbose = false
hcatalog.drop.and.create.table = false
incremental.last.value = 30000
db.connect.string = jdbc:mysql://quickstart.cloudera:3306/retail_db
codegen.output.delimiters.escape = 0
codegen.output.delimiters.enclose.required = false
codegen.input.delimiters.field = 0
mainframe.input.dataset.type = p
split.limit = null
hbase.create.table = false
db.require.password = true
hdfs.append.dir = true
db.table = orders_inc
codegen.input.delimiters.escape = 0
accumulo.create.table = false
import.fetch.size = null
codegen.input.delimiters.enclose.required = false
db.username = retail_dba
reset.onemapper = false
codegen.output.delimiters.record = 10
import.max.inline.lob.size = 16777216
sqoop.throwOnError = false
hbase.bulk.load.enabled = false
hcatalog.create.table = false
db.clear.staging.table = false
incremental.col = order_id
codegen.input.delimiters.record = 0
```

```
[cloudera@quickstart ~]$ sqoop job --show import_orders_inc | grep incremental
25/12/04 03:17:14 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.13.0
25/12/04 03:17:15 ERROR sqoop.SqoopOptions: It seems that you have launched a Sqoop metastore job via
25/12/04 03:17:15 ERROR sqoop.SqoopOptions: Oozie with sqoop.metastore.client.record.password disabled.
25/12/04 03:17:15 ERROR sqoop.SqoopOptions: But this configuration is not supported because Sqoop can't
25/12/04 03:17:15 ERROR sqoop.SqoopOptions: prompt the user to enter the password while being executed
25/12/04 03:17:15 ERROR sqoop.SqoopOptions: as Oozie tasks. Please enable sqoop.metastore.client.record
25/12/04 03:17:15 ERROR sqoop.SqoopOptions: .password in sqoop-site.xml, or provide the password
25/12/04 03:17:15 ERROR sqoop.SqoopOptions: explicitly using --password in the command tag of the Oozie
25/12/04 03:17:15 ERROR sqoop.SqoopOptions: workflow file.
incremental.last.value = 30000
incremental.col = order_id
incremental.mode = AppendRows
[cloudera@quickstart ~]$ cloudera@quickstart:~
```

4. Adding the remaining rows:

```
[cloudera@quickstart ~]$ mysql -uroot -pcloudera
Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 1634
Server version: 5.1.73 Source distribution

Copyright (c) 2000, 2013, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> use retail_db
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> insert into orders_inc select * from orders where order_id>30000 ;
Query OK, 38883 rows affected (0.14 sec)
Records: 38883 Duplicates: 0 Warnings: 0

mysql> commit;
Query OK, 0 rows affected (0.00 sec)
```

5. Re-executing the job:

```
25/12/04 03:18:48 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN(`order_id`), MAX(`order_id`) FROM `orders_in
c` WHERE ( `order_id` > 30000 AND `order_id` <= 68883 )
25/12/04 03:18:48 INFO db.IntegerSplitter: Split size: 9720; Num splits: 4 from: 30001 to: 68883
25/12/04 03:18:48 INFO mapreduce.JobSubmitter: number of splits:4
25/12/04 03:18:48 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1763910622083_0029
25/12/04 03:18:48 INFO impl.YarnClientImpl: Submitted application application_1763910622083_0029
25/12/04 03:18:49 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_17639106220
83_0029/
25/12/04 03:18:49 INFO mapreduce.Job: Running job: job_1763910622083_0029
25/12/04 03:18:55 INFO mapreduce.Job: Job job_1763910622083_0029 running in uber mode : false
25/12/04 03:18:55 INFO mapreduce.Job: map 0% reduce 0%
```

We can see that the request has limits 30000 to 68883.

6. Running job show again, we can see that the last value has changed from 30000 to 68883:

```
25/12/04 03:28:36 ERROR sqoop.SqoopOptions: workflow file.  
incremental.last.value = 68883  
incremental.col = order_id  
incremental.mode = AppendRows
```

Lab Hive

1. Importing the table “**stock_eod**” to Hive:

```
|hive> CREATE DATABASE stock;  
  
[cloudera@quickstart ~]$ sqoop import \  
> --connect "jdbc:mysql://quickstart.cloudera:3306"/nyse \  
> --username root \  
> --password cloudera \  
> --table stock_eod \  
> --hive-import \  
> --hive-database stock  
  
Failed Shuffles=0  
Merged Map outputs=0  
GC time elapsed (ms)=1432  
CPU time spent (ms)=10880  
Physical memory (bytes) snapshot=1517088768  
Virtual memory (bytes) snapshot=6312824832  
Total committed heap usage (bytes)=1421869056  
File Input Format Counters  
Bytes Read=0  
File Output Format Counters  
Bytes Written=25987047  
25/12/06 08:34:27 INFO mapreduce.ImportJobBase: Transferred 24.7832 MB in 34.0826 seconds (744.6024 KB/sec)  
25/12/06 08:34:27 INFO mapreduce.ImportJobBase: Retrieved 567649 records.  
25/12/06 08:34:27 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `stock_eod` AS t LIMIT 1  
25/12/06 08:34:27 WARN hive.TableDefWriter: Column openprice had to be cast to a less precise type in Hive  
25/12/06 08:34:27 WARN hive.TableDefWriter: Column highprice had to be cast to a less precise type in Hive  
25/12/06 08:34:27 WARN hive.TableDefWriter: Column lowprice had to be cast to a less precise type in Hive  
25/12/06 08:34:27 WARN hive.TableDefWriter: Column closeprice had to be cast to a less precise type in Hive  
25/12/06 08:34:27 INFO hive.HiveImport: Loading uploaded data into Hive  
Logging initialized using configuration in jar:file:/usr/lib/hive/lib/hive-common-1.1.0-cdh5.13.0.jar!/hive-log.properties  
Time taken: 2.029 seconds  
Loading data to table stock.stock_eod  
Table stock.stock_eod stats: [numFiles=4, totalSize=25987047]  
OK  
Time taken: 0.664 seconds
```

2. Performing 3 processings of my choice (on the table stock_eod):

```
import org.apache.spark.{SparkConf, SparkContext}
import org.apache.spark.sql.hive.HiveContext
import org.apache.spark.sql.functions._ // Required for agg, avg, max, col, etc.

object SparkHive {
  def main(args: Array[String]): Unit = {

    val conf = new SparkConf().setAppName("SparkHive").setMaster("yarn-client")
    val sc = new SparkContext(conf)
    val sqlContext = new org.apache.spark.sql.hive.HiveContext(sc)
    val myDF = sqlContext.sql("select * from stock.stock_eod")

    println("--- Schema of stock_eod ---")
    myDF.printSchema()

    // =====
    // Processing 1: Calculate Average Closing Price per Stock
    // =====
    println("--- Processing 1: Average Closing Price per Stock ---")
    val avgCloseDF = myDF.groupBy("stockticker")
      .agg(avg("closeprice").as("avg_close_price"))
      .orderBy(desc("avg_close_price"))

    avgCloseDF.show(10) // Show top 10 results

    // =====
    // Processing 2: Identify High Volatility Days (High - Low)
    // =====
    println("--- Processing 2: Highest Volatility Days ---")
    val volatilityDF = myDF.withColumn("price_spread", col("highprice") - col("lowprice"))
      .select("stockticker", "tradedate", "price_spread", "volume")
      .orderBy(desc("price_spread"))

    volatilityDF.show(10)

    // =====
    // Processing 3: Find Max Volume per Stock
    // =====
    println("--- Processing 3: Max Volume per Stock ---")
    val maxVolDF = myDF.groupBy("stockticker")
      .agg(max("volume").as("max_volume"))
      .orderBy(desc("max_volume"))

    maxVolDF.show(10)

    sc.stop()
  }
}
```

Executing the jar on the cluster using spark-submit:

```
[cloudera@quickstart lab_sqoop]$ spark-submit --class SparkHive \
> --master yarn \
> --deploy-mode client \
> stock_eod.jar
```

Outputs:

```
--- Schema of stock_eod ---
root
|-- stockticker: string (nullable = true)
|-- tradedate: string (nullable = true)
|-- openprice: double (nullable = true)
|-- highprice: double (nullable = true)
|-- lowprice: double (nullable = true)
|-- closeprice: double (nullable = true)
|-- volume: long (nullable = true)

--- Processing 1: Average Closing Price per Stock ---
+-----+-----+
|stockticker| avg_close_price|
+-----+-----+
| BRK.A| 93787.53946360154|
| WFC-L| 744.0917692307694|
| BAC-L| 706.1042528735628|
| NVR| 547.6852873563215|
| MKL| 301.48463601532586|
| Y| 265.5966666666667|
| NBG| 256.56130268199234|
| WTM| 256.0481609195403|
| ALX| 246.60980842911866|
| BLK| 172.5124904214559|
+-----+-----+
only showing top 10 rows

--- Processing 2: Highest Volatility Days ---
+-----+-----+-----+-----+
|stockticker| tradedate| price_spread|volume|
+-----+-----+-----+-----+
| BRK.A| 10-Mar-2009| 11634.0| 3000|
| BRK.A| 07-Jan-2009| 6600.0| 1400|
| BRK.A| 24-Feb-2009| 5800.0| 3000|
| BRK.A| 10-Aug-2009| 5799.0| 1100|
| BRK.A| 02-Mar-2009| 5750.0| 2400|
| BRK.A| 05-Aug-2009| 5728.0| 1500|
| BRK.A| 08-Jan-2009| 5700.0| 1300|
| BRK.A| 12-Mar-2009| 5619.990000000005| 1600|
| BRK.A| 20-Jan-2009| 5524.990000000005| 1600|
| BRK.A| 23-Feb-2009| 5100.0| 3170|
+-----+-----+-----+-----+
only showing top 10 rows

--- Processing 3: Max Volume per Stock ---
+-----+-----+
|stockticker|max_volume|
+-----+-----+
| BAC| 1226792000|
| GE| 752731008|
| AIG| 583709400|
| RF| 570844000|
| CIT| 566983700|
| F| 541015000|
| WFC| 478614300|
| GM| 340927000|
| S| 244426200|
| AA| 242126000|
+-----+-----+
only showing top 10 rows
```